



A Julia Package for HPC Meta-Programming and Performance Portability

Philip Fackler


*In collaboration with Pedro Valero-Lara,
William Godoy, Het Mankad, Keita Teranishi,
and Jeffrey Vetter*



ORNL IS MANAGED BY UT-BATTELLE LLC
FOR THE US DEPARTMENT OF ENERGY



Why Julia?




- Easy to use
- Easy performance
- Easily supported (JIT on top of LLVM)
- Sensible syntax for science
- Easy for HPC packages to provide native syntax (not just wrapping)
- Integrated packaging and testing (reproducibility)
- Easy portability
 - (usually just works on a new system...from a developer's standpoint)

Motivation:

- Varied systems
- Varied user expertise

Programming Productivity Performance Portability Performance “Practicability”?




1
oneAPI

SYCL™



NVIDIA.
CUDA.

OpenMP®



JACC.jl, Julia for ACCELERATORS


Simple to use performance portable model

- High-level performance portability model
(parallel_for & parallel_reduce)
- Unified Julia front-end on top of multiple backends
 - Threads, CUDA, AMDGPU, oneAPI (Metal coming soon)
- Hide low-level, device-specific implementation
 - Memory layout, granularity, etc.
- ***Brings HPC closer to domain scientists***



<https://github.com/JuliaORNL/JACC.jl>

<https://ieeexplore.ieee.org/document/10820713>



Getting started with JACC



1. Add JACC as a dependency

```
julia> Pkg.add("JACC")
```

2. Set backend

```
julia> using JACC
```

```
julia> JACC.set_backend("cuda")
```

- Adds backend package

3. Restart julia

4. Load extension before use

```
julia> using JACC
```

```
julia> JACC.@init_backend
```

- Imports backend package -> activates extension

JACC paradigm basics - arrays



No JACC-defined array type

`JACC.array(x::Base.Array)`

Constructs a backend-specific array (copying to device if necessary)

`JACC.array_type()`

Provides type of array that would be returned by `JACC.array`, e.g., `CuArray`

```
function foo(x::JACC.array_type()) {Float64, 1})  
    # ...  
end
```

`JACC.ones`, `JACC.zeros`, `JACC.fill`

JACC paradigm basics - parallel_for



Moving from a serial loop to a JACC.parallel_for

```
for i = 1:N  
    @inbounds x[i] += alpha * y[i]  
end
```



```
JACC.parallel_for(N, alpha, dx, dy) do i, alpha, x, y  
    @inbounds x[i] += alpha * y[i]  
end
```

(or, using predefined function)

```
function axpy(i, alpha, x, y)  
    @inbounds x[i] += alpha * y[i]  
end  
JACC.parallel_for(SIZE, axpy, alpha, dx, dy)
```

JACC paradigm basics - parallel_reduce



Examples:

```
elem_sum = JACC.parallel_reduce (a)
elem_min = JACC.parallel_reduce (min, a)

dp = JACC.parallel_reduce (10, a, b) do i, a, b
    @inbounds a[i] * b[i]
end
```

API overview:

```
parallel_reduce(N, f, x...; op, init) -> typeof(init)
parallel_reduce(N, f, x...) = parallel_reduce(N, f, x...; op = +, init = 0.0)
parallel_reduce([op = +,] a::AbstractArray; init = default_init(eltype(a), op)) -> typeof(init)
# op ∈ {+, *, min, max, custom}
```

Extra Things



LaunchSpec for fine granularity

```
launch_spec(; kw...)  
  
accepts keyword arguments:  
    stream  
    threads  
    blocks  
    shmem_size  
    sync  
  
spec = JACC.launch_spec(;  
    sync = false, threads = 1000)  
  
JACC.parallel_for(spec, N, a_device) do i, a  
    @inbounds a[i] += 5.0  
end
```

JACC.shared: on-chip shmem

```
function spectral(i, j, image, filter,  
    num_bands)  
    for b in 1:num_bands  
        @inbounds image[b, i, j] *= filter[j]  
    end  
end  
  
(just add one line)  
  
function spectral(i, j, image, filter,  
    num_bands)  
    filter_sh = JACC.shared(filter)  
    for b in 1:num_bands  
        @inbounds image[b, i, j] *= filter_sh[j]  
    end  
end
```

Extra Things (multi-device nodes) (experimental)



JACC.Multi: multi-device “parallelism”

```
alpha = 2.5
dx = JACC.Multi.array (x)
dy = JACC.Multi.array (y)

JACC.Multi.parallel_for (N,a,dx,dy) do i, a, x, y
    @inbounds x[i] += alpha * y[i]
end
```

With extra methods for handling ghost elements

JACC.Async: multi-device “concurrency”

```
JACC.Async.parallel_for (1, ...)
JACC.Async.parallel_for (2, ...)
JACC.Async.synchronize ()

# parallel_reduce return value is on device
res_d = JACC.Async.parallel_reduce (1, ...)
res = JACC.to_host (res_d) []
```

Applications & Future Work

- Ongoing Efforts
 - Performance benchmarking (closing gaps)
 - More intuitive kernel launch
 - JACC.Auto : Autotuning
 - Have ideas? Send us a message!
- Examples and applications using JACC
 - Look at the [tests](#)
 - Tatiana's [7-point stencil](#)
 - <https://github.com/JuliaORNL/MiniVATES.jl>
 - <https://github.com/JuliaORNL/JACC-applications>
 - <https://github.com/JuliaORNL/GrayScott.jl>
- Other HPC Julia tutorial resources
 - [SC24 tutorial](#)

Project:

- William F Godoy (ORNL)
- Pedro Valero Lara (ORNL)
- Philip W Fackler (ORNL)
- Narasinga Rao Miniskar (ORNL)
- Aaron Young (ORNL)
- Keita Teranishi (ORNL)

Contributions:

- Het Mankad (ORNL)
- Julian Samaroo (MIT)
- Michel Schanen (ANL)
- Karl Pierce (Flatiron Institute)
- Kelly Tang
- Tatiana Melnichenko

Thanks!

This research used resources of the Oak Ridge Leadership Computing Facility and the Experimental Computing Laboratory at the Oak Ridge National Laboratory, which is supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.

Sponsors:

ASCR Competitive Portfolios: Fairbanks project
ASCR NGSST PESO and S4PST projects
ASCR Facilities: OLCF/NERSC

Questions?